

Lab1 - IE406

Sharvil Sheth - 201901201, Prabhav Shah - 201901216, Mohil Desai - 201901301, Dhaval Vaidya - 201901462, Vishvesh Patel - 201901464

Importing Libraries

```
In [45]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
from mpl_toolkits.mplot3d import Axes3D
import math
from sklearn import linear_model
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression
```

Q1:

```
In [46]: #read data from the CSV
data = pd.read_excel('Data for Lab 1.xlsx')
data.head()
```


```
Out[46]:
```

	x	y
0	3504	18.0
1	3693	15.0
2	3436	18.0
3	3433	16.0
4	3449	17.0

```
In [47]: #store X and Y
x = data['x']
y = data['y']
```

```
In [48]: theta = np.linspace(-10,10,201)
lofTheta = np.square(theta)

#plot
plt.plot(theta, lofTheta)
plt.xlabel('theta')
plt.ylabel('L(theta)')
plt.show()
```



Observations:

Here we observe that the curve of $L(\theta)$ is parabolic in nature. It attains the minimum value at $L(\theta) = 0$ at $\theta = 0$.

Q2:


In [49]:

```
theta_1 = np.copy(theta)
theta_2 = np.copy(theta)

lOftheta = np.zeros((201, 201))
for i in range(201):
    for j in range(201):
        lOftheta[i][j] = ((theta_1[i]*theta_1[i]) + (theta_2[j]*theta_2[j]))

theta_1, theta_2 = np.meshgrid(theta_1, theta_2)

#plot
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(X = theta_1, Y = theta_2, Z = lOftheta)
ax.set_xlabel('theta_1')
ax.set_ylabel('theta_2')
ax.set_zlabel('L(theta)')
plt.show()
```



Here we get a contour plot of $L(\theta)$ for different values of θ_1 and θ_2 . $L(\theta)$ attains its minimum value $L(\theta) = 0$ at $(\theta_1, \theta_2) = (0, 0)$

Q3a:

In [50]:

```
x = np.array(data['x'])
y = np.array(data['y'])

ones = np.ones(len(x))
ones = ones.reshape((94,1))

x = x.reshape((94,1))
x = np.append(ones, x,axis = 1)

iterations = 50
alpha = 0.0000001
th_0 = 50
th_1 = 0
n = len(x)
final_theta_0 = []
final_theta_1 = []
final_costs = []
for i in range(iterations):

    y_predicted = th_1 * x[:,1] + th_0 * x[:,0]

    cost = sum((y-y_predicted)**2)

    delta_th_1 = -(2/n)*sum(x[:,1]*(y-y_predicted))
    delta_th_0 = -(2/n)*sum(y-y_predicted)


    th_1 = th_1 - alpha * delta_th_1
    th_0 = th_0 - alpha * delta_th_0

    final_theta_0.append(th_0)
    final_theta_1.append(th_1)

    final_costs.append(cost)

final_theta_0 = np.array(final_theta_0)
final_theta_1 = np.array(final_theta_1)
final_costs = np.array(final_costs)

#plot
fig = plt.figure(figsize=(10, 10))
ax = plt.axes(projection='3d')
ax.plot3D(final_theta_0,final_theta_1, final_costs)
ax.set_xlabel('theta_0')
ax.set_ylabel('theta_1')
ax.set_zlabel('cost')
plt.show()
```



```
In [51]: min_L_theta = min(final_costs) # minimum value value of L(theta)
min_theta_0 = final_theta_0[np.argmin(final_costs)]
min_theta_1 = final_theta_1[np.argmin(final_costs)]
print(min_theta_0,min_theta_1,min_L_theta)
```

```
49.999996705417516 -0.008833886263553227 1576.7046182979047
```

Q3b:

```
In [52]: x_temp = x
x_temp = x_temp.reshape((94,2))
y = y.reshape((94,1))
temp_x_t = x_temp.T.dot(x_temp)
theta = np.linalg.pinv(temp_x_t).dot(x_temp.T).dot(y)
print(theta)
```

```
[[ 4.92376299e+01]
 [-8.61193478e-03]]
```

Q4:

```
In [53]: y_predicted = theta[0]*x[:,0] + theta[1]*x[:,1]
y = y.reshape(len(y_predicted))
L_theta = sum((y - y_predicted)**2)

print(L_theta)

theta_rand = np.array([55,-0.005])
y_predicted = theta_rand[0]*x[:,0] + theta_rand[1]*x[:,1]
y = y.reshape(len(y_predicted))
L_theta_rand = sum((y - y_predicted)**2)

print(L_theta_rand)

1572.6503668922924
27837.063900000001
```

Q5:

```
In [54]: # Creating our dataset
X = np.array([[1, 2],[2, 4],[3, 6],[4, 8]])
Y = np.array([2,3,4,5])
Y = Y.reshape((4,1))
```

(a) using scikit-learn library

```
In [55]: #Creating model() instance and training our model on data (X,Y)
model = linear_model.LinearRegression()
model.fit(X,Y)
```

```
Out[55]: LinearRegression()
```

```
In [56]: model.score(X,Y)
```

```
Out[56]: 1.0
```

```
In [57]: model.coef_
```

```
Out[57]: array([[0.2, 0.4]])
```

```
In [58]: model.intercept_
```

```
Out[58]: array([1.])
```

We get coefficients value [theta1,theta2] as [0.2,0.4] and intercept value as 1.

In equation, $y = 0.2x_1 + 0.4x_2 + 1$

```
In [59]: #prediction with the help of trained model
model.predict([[5,10], [6,12], [7,14]])
```

```
Out[59]: array([[6.],
               [7.],
               [8.]])
```

(b) Using normal equations

```
In [60]: def normal_equation(X,y):
          x_tmp = X.T.dot(X)
          theta = np.linalg.inv(x_tmp).dot(X.T).dot(y)
          return theta

          XX = np.c_[np.ones((4,1)), X]
          theta = normal_equation(XX, Y)
```

```
-----
LinAlgError                                Traceback (most recent call last)
<ipython-input-60-ab9013afbd25> in <module>
      5
      6 XX = np.c_[np.ones((4,1)), X]
----> 7 theta = normal_equation(XX, Y)

<ipython-input-60-ab9013afbd25> in normal_equation(X, y)
      1 def normal_equation(X,y):
      2     x_tmp = X.T.dot(X)
----> 3     theta = np.linalg.inv(x_tmp).dot(X.T).dot(y)
      4     return theta
      5

~\Anaconda3\lib\site-packages\numpy\linalg\linalg.py in inv(a)
    549     signature = 'D->D' if isComplexType(t) else 'd->d'
    550     extobj = get_linalg_error_extobj(_raise_linalgerror_singular)
--> 551     ainv = _umath_linalg.inv(a, signature=signature, extobj=extobj)
    552     return wrap(ainv.astype(result_t, copy=False))
    553

~\Anaconda3\lib\site-packages\numpy\linalg\linalg.py in _raise_linalgerror_
singular(err, flag)
     95
     96 def _raise_linalgerror_singular(err, flag):
--> 97     raise LinAlgError("Singular matrix")
     98
     99 def _raise_linalgerror_nonposdef(err, flag):

LinAlgError: Singular matrix
```

We got Singular Matrix Error because $\det(X'X) = 0$. Even though, scikit learn can solve the model as it uses pseudo inverse instead of normal inverse. So, if we use pinv function instead of inv function, we can still get the correct results.

Q6a:

```
In [61]: dataset = pd.read_excel(r'Real estate valuation data set.xlsx')
          X = dataset.iloc[:,1:7]
          Y = dataset.iloc[:,7]
```

```
In [62]: X_train, X_test, Y_train, Y_test = train_test_split( X, Y, test_size = 1/3,
```

```
In [63]: reg = LinearRegression()
reg = reg.fit(X_train, Y_train)

Y_pred = reg.predict(X_test)
```

```
In [64]: rms = math.sqrt(mean_squared_error(Y_test, Y_pred))
print(rms)
```

8.743764087799477

```
In [65]: print(reg.coef_)
print(reg.score(X_test, Y_test))

[ 5.43155901e+00 -3.16000575e-01 -4.33141896e-03  1.21738116e+00
 2.52174349e+02 -4.95885551e+00]
0.5417678968095335
```

Q6b:

The value of the regression coefficients does not anyhow show the importance of each different feature. So, it will be entirely wrong to assume that larger coefficients mean more important features.

Q6c:

```
In [66]: min_max = MinMaxScaler()
X_normalised = min_max.fit_transform(X)
```

```
In [67]: X_n_train, X_n_test, Y_n_train, Y_n_test = train_test_split(X_normalised,
reg_n = LinearRegression()
reg_n = reg_n.fit(X_n_train, Y_n_train)

Y_n_pred = reg_n.predict(X_n_test)
```

```
In [68]: rms_norm = math.sqrt(mean_squared_error(Y_n_test, Y_n_pred))
print(rms_norm)
```

8.743764087799743

```
In [69]: print(reg_n.coef_)
print(reg_n.score(X_n_test, Y_n_test))


[ 4.97892873 -13.84082518 -28.00105628  12.17381156  20.80942729
 -0.45988426]
0.5417678968095057
```

Q6d:

Gaussian

```
In [70]: residual = Y_n_pred - Y_n_test
residual.hist(bins=30,range = [-30,30])
```

```
Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x24349de02c8>
```



Q6e:

```
In [71]: dataset.corr()
```

```
Out[71]:
```

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitu
No	1.000000	-0.048634	-0.032808	-0.013573	-0.012699	-0.010110	-0.0110
X1 transaction date	-0.048634	1.000000	0.017542	0.060880	0.009544	0.035016	-0.0410
X2 house age	-0.032808	0.017542	1.000000	0.025622	0.049593	0.054420	-0.0485
X3 distance to the nearest MRT station	-0.013573	0.060880	0.025622	1.000000	-0.602519	-0.591067	-0.8063
X4 number of convenience stores	-0.012699	0.009544	0.049593	-0.602519	1.000000	0.444143	0.4490
X5 latitude	-0.010110	0.035016	0.054420	-0.591067	0.444143	1.000000	0.4129
X6 longitude	-0.011059	-0.041065	-0.048520	-0.806317	0.449099	0.412924	1.0000
Y house price of unit area	-0.028587	0.087529	-0.210567	-0.673613	0.571005	0.546307	0.5232

In [72]:

```
X_optimum = dataset.iloc[:,[3,4,5,6]]  
  
min_max = MinMaxScaler()  
X_opt_normalised = min_max.fit_transform(X_optimum)
```

In [73]:

```
X_opt_train, X_opt_test, Y_opt_train, Y_opt_test = train_test_split( X_opt_   
  
reg_opt = LinearRegression()  
reg_opt = reg.fit(X_opt_train, Y_opt_train)  
  
Y_opt_pred = reg_opt.predict(X_opt_test)
```

In [74]:

```
rms_opt = math.sqrt(mean_squared_error(Y_opt_test, Y_opt_pred))  
print(rms_opt)
```

9.000203768125342

In [75]:

```
print(reg_opt.coef_)  
print(reg_opt.score(X_opt_test, Y_opt_test))
```

[-29.39425081 11.30570064 20.71949247 -0.59106937]
0.5144954176958135

In [76]:

```
residual = Y_opt_pred - Y_opt_test  
residual.hist(bins=30, range = [-30,30])
```

Out[76]: <matplotlib.axes._subplots.AxesSubplot at 0x2434ae34e88>

