

IE406 Machine Learning

Lab Assignment - 7

Group 14

201901466: Miti Purohit

202001430: Aryan Shah

#Question 1

Cars24 is the most popular website of used vehicles for sale, yet it's very difficult to collect all of them in the same place. Among all cities, data from 5 major cities which include Hyderabad, New Delhi, Mumbai, Bangalore, and Chennai is collected. Develop an algorithm for predict price of car.

Data link: <https://www.kaggle.com/balajimummidi/used-cars-in-cars24>

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import mean_squared_error
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
from sklearn.linear_model import SGDRegressor
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, Dense
```

```
[ ]: df = pd.read_csv('Cars24.csv', index_col=0)
df.head()
```

```
[ ]: Car Brand      Model  Price  Model Year  Location  Fuel \
0   Hyundai      EonERA PLUS  330399      2016  Hyderabad  Petrol
1   Maruti      Wagon R 1.0LXI  350199      2011  Hyderabad  Petrol
2   Maruti      Alto K10LXI  229199      2011  Hyderabad  Petrol
3   Maruti      RitzVXI BS IV  306399      2011  Hyderabad  Petrol
4   Tata      NanoTWIST XTA  208699      2015  Hyderabad  Petrol

Driven (Kms)      Gear  Ownership  EMI (monthly)
0          10674  Manual          2           7350
1          20979  Manual          1           7790
```

2	47330	Manual	2	5098
3	19662	Manual	1	6816
4	11256	Automatic	1	4642

```
[ ]: df.dropna(inplace=True)
df
```

```
[ ]:
Car Brand      Model  Price  Model Year  Location \
0    Hyundai      EonERA PLUS  330399      2016  Hyderabad
1    Maruti      Wagon R 1.0LXI  350199      2011  Hyderabad
2    Maruti      Alto K10LXI  229199      2011  Hyderabad
3    Maruti      RitzVXI BS IV  306399      2011  Hyderabad
4    Tata      NanoTWIST XTA  208699      2015  Hyderabad
...    ...      ...      ...      ...      ...
5913  Toyota      Fortuner3.0 AT 4X2  1234899      2012  Chennai
5914  Toyota      Innova2.5 VX 8 STR BS IV  892699      2012  Chennai
5915  Maruti      Wagon R 1.0VXI  381399      2014  Chennai
5916  Hyundai      i10SPORTZ 1.2 KAPPA2 0  310899      2011  Chennai
5917  Maruti      Wagon R DuoLXI LPG  159999      2007  Chennai
```

	Fuel	Driven (Kms)	Gear	Ownership	EMI (monthly)
0	Petrol	10674	Manual	2	7350
1	Petrol	20979	Manual	1	7790
2	Petrol	47330	Manual	2	5098
3	Petrol	19662	Manual	1	6816
4	Petrol	11256	Automatic	1	4642
...
5913	Diesel	197177	Automatic	1	27470
5914	Diesel	115553	Manual	2	19858
5915	Petrol	24663	Manual	1	8484
5916	Petrol	30159	Manual	1	6916
5917	Petrol + LPG	51247	Manual	2	3559

[5653 rows x 10 columns]

```
[ ]: def getVecForm(vocab, df):
      for i in range(len(vocab)):
          df.replace(vocab[i], i, inplace=True)
      return df
```

```
[ ]: df['Fuel'] = getVecForm(np.unique(df['Fuel']), df['Fuel'])
df['Car Brand'] = getVecForm(np.unique(df['Car Brand']), df['Car Brand'])
df['Model'] = getVecForm(np.unique(df['Model'].astype(str)), df['Model'])
df['Location'] = getVecForm(np.unique(df['Location']), df['Location'])
df['Gear'] = getVecForm(np.unique(df['Gear'].astype(str)), df['Gear'])
df
```

```
[ ]:      Car Brand  Model    Price  Model Year  Location  Fuel  Driven (Kms)  \
0          7    308   330399        2016         3    2      10674
1         15    811   350199        2011         3    2      20979
2         15    31   229199        2011         3    2      47330
3         15   611   306399        2011         3    2      19662
4         22   526   208699        2015         3    2      11256
...      ...    ...    ...      ...      ...    ...    ...
5913       23   373  1234899        2012         1    0     197177
5914       23   453   892699        2012         1    0     115553
5915       15   813   381399        2014         1    2      24663
5916        7   884   310899        2011         1    2      30159
5917       15   820   159999        2007         1    4      51247
```

```
      Gear  Ownership  EMI (monthly)
0         1         2         7350
1         1         1         7790
2         1         2         5098
3         1         1         6816
4         0         1         4642
...      ...      ...      ...
5913      0         1        27470
5914      1         2        19858
5915      1         1         8484
5916      1         1         6916
5917      1         2         3559
```

[5653 rows x 10 columns]

```
[ ]: X = df.drop(['Price'], axis=1).to_numpy()
y = df['Price'].to_numpy()

scaler = StandardScaler()
X = scaler.fit_transform(X)
```

Logistic Regression

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
→random_state=42)
```

```
[ ]: reg = LogisticRegression(random_state=0, solver='liblinear', max_iter=1000).
→fit(X_train, y_train)
y_pred = reg.predict(X_test)
```


```
[ ]: reg.score(X_test, y_test)
```

```
[ ]: 0.0008841732979664014
```


```
[ ]: print(mean_squared_error(y_test, y_pred, squared=False))
```

223357.67037352617

```
[ ]: plt.figure(figsize=[10,8])  
plt.plot(y_test, 'b', alpha=0.5, label='Actual')  
plt.plot(y_pred, 'k', alpha=0.5, label='Predicted')  
plt.legend()  
plt.show()
```




```
[ ]: plt.figure(figsize=[8,8])  
plt.scatter(y_test, y_pred, color='b')  
plt.xlabel('True Values [Price]')  
plt.ylabel('Predictions [Price]')  
plt.show()
```



```
[ ]: error = abs(y_test-y_pred)
```


```
[ ]: plt.figure(figsize=[8,8])
plt.hist(error, bins = 10, color='b', alpha=0.5)
plt.xlabel("Prediction Error [Price]")
plt.ylabel("Count")
plt.show()
```



PCA followed by Logistic Regression

```
[ ]: pca = PCA(n_components=X.shape[1])
pca.fit_transform(X)
ex_var_ratio = pca.explained_variance_ratio_

plt.figure(figsize=[8,8])
plt.plot(ex_var_ratio, color='b', linestyle='-', marker='x')
plt.show()
```



```
[ ]: pca = PCA(n_components=3)
X_new = pca.fit_transform(X)
print(np.sum(pca.explained_variance_ratio_))
```

0.5177464662742188

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size=0.2,
→random_state=42)
```

```
[ ]: reg = LogisticRegression(max_iter=1000).fit(X_train, y_train)
y_pred = reg.predict(X_test)
```


```
[ ]: reg.score(X_test, y_test)
```

[]: 0.0008841732979664014


```
[ ]: print(mean_squared_error(y_test, y_pred, squared=False))
```

189427.8861986012

```
[ ]: plt.figure(figsize=[8,6])
plt.plot(y_test, 'b', alpha=0.5, label='Actual')
plt.plot(y_pred, 'k', alpha=0.5, label='Predicted')
plt.legend()
plt.show()
```




```
[ ]: plt.figure(figsize=[8,8])
plt.scatter(y_test, y_pred, color='b')
plt.xlabel('True Values [Price]')
plt.ylabel('Predictions [Price]')
plt.show()
```

```
[ ]: error = abs(y_test-y_pred)
```

```
[ ]: plt.figure(figsize=[8,8])
plt.hist(error, bins = 10, color='b', alpha=0.5)
plt.xlabel("Prediction Error [Price]")
plt.ylabel("Count")
plt.show()
```



Stochastic Gradient Descent

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
n_samples, n_features = X_train.shape[0], X_train.shape[1]
reg = make_pipeline(StandardScaler(),
                    SGDRegressor(max_iter=1000, tol=1e-3))
reg.fit(X_train, y_train)
```


```
[ ]: Pipeline(steps=[('standardscaler', StandardScaler()),
                    ('sgdregressor', SGDRegressor())])
```

```
[ ]: y_pred = reg.predict(X_test)
```


```
[ ]: reg.score(X_test, y_test)
```

```
[ ]: 0.9999999806600117
```

```
[ ]: plt.figure(figsize=[10,8])  
plt.plot(y_test, 'b', alpha=0.5, label='Actual')  
plt.plot(y_pred, 'k', alpha=0.5, label='Predicted')  
plt.legend()  
plt.show()
```




```
[ ]: plt.figure(figsize=[8,8])  
plt.scatter(y_test, y_pred, color='b')  
plt.xlabel('True Values [Price]')  
plt.ylabel('Predictions [Price]')  
plt.show()
```



```
[ ]: error = abs(y_test-y_pred)
```

```
[ ]: plt.figure(figsize=[8,8])  
plt.hist(error, bins = 10, color='b', alpha=0.5)  
plt.xlabel("Prediction Error [Price]")  
plt.ylabel("Count")  
plt.show()
```



###Result

The graphs are shown above.

###Observation/Justification

The error in prediction of car price is highest for Logistic regression model with and without PCA. Stochastic GD and Neural Network models perform much better than original Logistic Regression model. This is clearly evident from the linear line, the slope of which is almost 1

#Question 2

Fashion-MNIST is a dataset of Zalando's article images—consisting of a training set of 60,000 examples and a test set of 10,000 examples. Each example is a 28x28 grayscale image, associated with a label from 10 classes. That means one image contains 784 pixel and pixel-value is an integer between 0 and 255. Make a classification model to classify the product.

Data link: <https://www.kaggle.com/zalando-research/fashionmnist>

#Answer

###Code

```
[ ]: import numpy as np
import pandas as pd
import time
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objects as go
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.metrics import confusion_matrix, \
    →ConfusionMatrixDisplay, classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, recall_score, precision_score, \
    →f1_score
from sklearn import svm
import tensorflow as tf
import keras as keras
```

```
[ ]: training_images = pd.read_csv('fashion.csv')
training_labels = training_images['label']
training_images = training_images[training_images.columns[1:]]
test_images = pd.read_csv('fashion.csv')
test_labels = test_images['label']
test_images = test_images[test_images.columns[1:]]
training_images=StandardScaler().fit_transform(training_images)
test_images=StandardScaler().fit_transform(test_images)
```

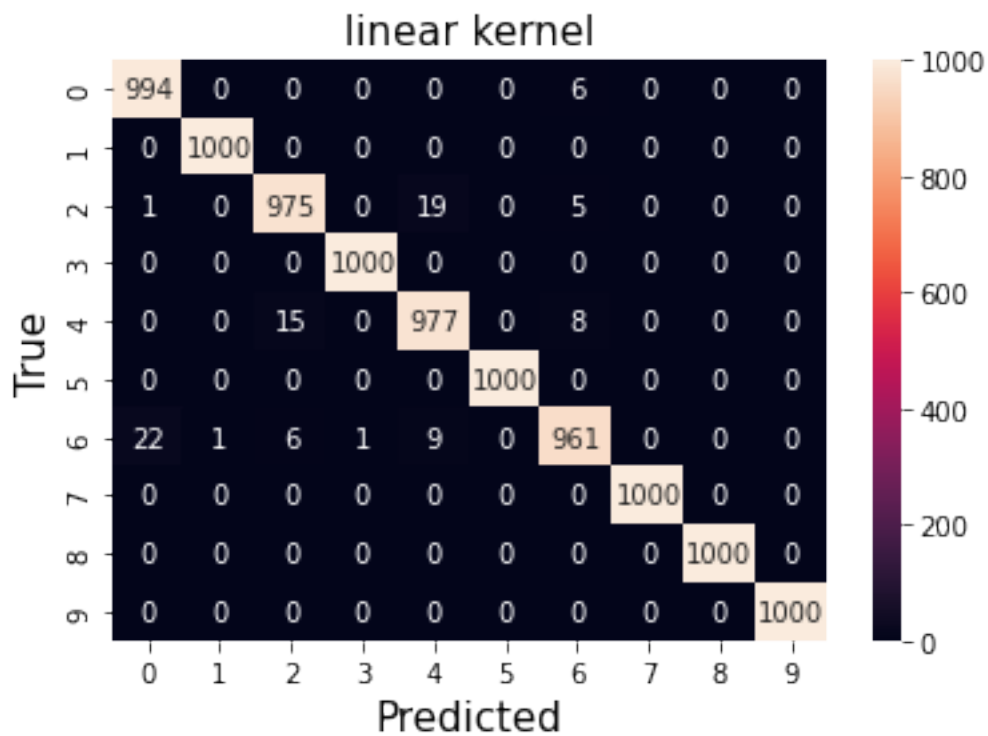
SVM

```
[ ]: #Plot confusion matrix
def plot_confusion_matrix(title,y_test,y_pred):

    confusionMatrix = confusion_matrix(y_test,y_pred)
    sns.heatmap(confusionMatrix,annot=True, fmt='.4g')
    plt.ylabel('True',fontsize=15)
    plt.xlabel('Predicted',fontsize=15)
    plt.title(title,fontsize=15)
    plt.show()
    print(classification_report(y_test,y_pred))
```


```
[ ]: # SVM
def svm_func(x_train, x_test, y_train, y_test):
    k = ['linear', 'poly']
    for i in range(len(k)):
        clf = svm.SVC(kernel=k[i])
        clf.fit(x_train, y_train)
        y_pred = clf.predict(x_test)
        plot_confusion_matrix(k[i]+' kernel', y_test, y_pred)

svm_func(training_images, test_images, training_labels, test_labels)
```



	precision	recall	f1-score	support
0	0.98	0.99	0.99	1000
1	1.00	1.00	1.00	1000
2	0.98	0.97	0.98	1000
3	1.00	1.00	1.00	1000
4	0.97	0.98	0.97	1000
5	1.00	1.00	1.00	1000
6	0.98	0.96	0.97	1000
7	1.00	1.00	1.00	1000
8	1.00	1.00	1.00	1000
9	1.00	1.00	1.00	1000

accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000




	precision	recall	f1-score	support
0	0.89	0.88	0.88	1000
1	1.00	0.96	0.98	1000
2	0.88	0.78	0.83	1000
3	0.90	0.90	0.90	1000
4	0.88	0.75	0.81	1000
5	0.97	0.96	0.97	1000
6	0.62	0.82	0.71	1000
7	0.93	0.97	0.95	1000
8	1.00	0.97	0.99	1000
9	0.98	0.95	0.97	1000

accuracy			0.90	10000
macro avg	0.90	0.90	0.90	10000
weighted avg	0.90	0.90	0.90	10000

Logistic Regression

```
[ ]: clf = LogisticRegression(max_iter=10000).fit(training_images, training_labels)
accuracy_inbuilt = accuracy_score(test_labels, clf.predict(test_images))
pred_data = clf.predict(test_images)
```

```
[ ]: cm = confusion_matrix(test_labels, pred_data)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
fig, ax = plt.subplots(figsize=(10,10))
disp.plot(ax=ax)
plt.show()
print(classification_report(test_labels, pred_data))
```




precision recall f1-score support

0	0.94	0.95	0.95	1000
1	1.00	1.00	1.00	1000
2	0.92	0.92	0.92	1000
3	0.97	0.99	0.98	1000
4	0.92	0.93	0.92	1000
5	1.00	1.00	1.00	1000
6	0.90	0.87	0.89	1000
7	0.99	1.00	1.00	1000
8	1.00	1.00	1.00	1000
9	1.00	1.00	1.00	1000
accuracy			0.96	10000
macro avg	0.96	0.96	0.96	10000
weighted avg	0.96	0.96	0.96	10000

Linear Discriminant Analysis

```
[ ]: clf = LinearDiscriminantAnalysis()
      clf.fit(training_images, training_labels)
      pred_data = clf.predict(test_images)

[ ]: cm = confusion_matrix(test_labels, pred_data)
      disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
      fig, ax = plt.subplots(figsize=(10,10))
      disp.plot(ax=ax)
      plt.show()
      print(classification_report(test_labels, pred_data))
```



	precision	recall	f1-score	support
0	0.83	0.82	0.83	1000
1	1.00	0.95	0.97	1000
2	0.79	0.76	0.77	1000
3	0.87	0.89	0.88	1000
4	0.78	0.82	0.80	1000
5	0.90	0.93	0.92	1000
6	0.64	0.66	0.65	1000
7	0.90	0.90	0.90	1000
8	0.98	0.95	0.96	1000
9	0.94	0.93	0.93	1000
accuracy			0.86	10000

macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

###Result

The plots and graphs for SVM, Logistic Regression and Linear Discriminant Analysis is as shown above.

###Observation/Justification

The SVM linear kernel model performs the best among all the models chosen.